# 274-[DF]-Lab - Introduction to Amazon Aurora

### Introduction to Amazon Aurora

### Overview

This lab introduces you to Amazon Aurora and provides you with a basic understanding of how to use Aurora. You will follow the steps to create an Aurora instance and then connect to it.

### **Topics covered**

After completing this lab, you will be able to:

- Create an Aurora instance
- Connect to a pre-created Amazon Elastic Compute Cloud (Amazon EC2) instance
- · Configure the Amazon EC2 instance to connect to Aurora
- · Query the Aurora instance

### **Duration**

This lab requires approximately 40 minutes to complete.

### Icon key

Various icons are used throughout this lab to call attention to different types of instructions and notes. The following list explains the purpose for each icon:

- E Command: A command that you must run.
- Expected output: A sample output that you can use to verify the output of a command or edited file.
- F Note: A hint, tip, or important guidance.
- ① CAUTION: Information of special interest or importance (not so important to cause problems with the equipment or data if you miss it, but it could result in the need to repeat certain steps).
- Consider: A moment to pause to consider how you might apply a concept in your own environment or to initiate a conversation about the topic at hand.
- Copy edit: A time when copying a command, script, or other text to a text editor (to edit specific variables within it) might be easier than editing directly in the command line or terminal.
- 🖒 Task complete: A conclusion or summary point in the lab.

### **AWS** service restrictions

In this lab environment, access to AWS services and service actions might be restricted to the ones that you need to complete the lab instructions. You might encounter errors if you attempt to access other services or perform actions beyond the ones that this lab describes

### **Prerequisites**

To successfully complete this lab, you should have some experience using the Linux operating system and have a basic understanding of structured query language (SQL).

### Accessing the AWS Management Console

- 1. At the upper-right corner of these instructions, choose > Start Lab.
  - F Note: If you get an Access Denied error, close the error box, and choose ▶ Start Lab again.
- 2. The lab status can be interpreted as follows:
  - A red circle next to <u>AWS</u> at the upper-left corner of this page indicates the lab has not been started.
  - · A yellow circle next to AWS at the upper-left corner of this page indicates the lab is starting.
  - A green circle next to AWS at the upper-left corner of this page indicates the lab is ready.

Wait for the lab to be ready before proceeding.

3. At the top of these instructions, choose the green circle next to AWS.

This option opens the AWS Management Console in a new browser tab. The system automatically sign you in.

- Note: If a new browser tab does not open, a banner or icon at the top of your browser indicates that your browser is preventing the site from opening pop-up windows. Choose the banner or icon, and choose Allow pop-ups.
- 4. Arrange the AWS Management Console tab so that it displays along side these instructions. Ideally, you should be able to see both browser tabs at the same time so that you can follow the lab steps.
  - 1 Caution: Do not change the lab Region unless specifically instructed to do so.

It takes a few minutes to provision the resources necessary to complete this lab.

### Introducing the technologies

### Amazon Aurora

Aurora is a fully managed, MySQL-compatible, relational database engine that combines the performance and reliability of high-end commercial databases with the simplicity and cost-effectiveness of open-source databases. It delivers up to five times the performance of MySQL without requiring changes to most of your existing applications that use MySQL databases.

#### Amazon Elastic Compute Cloud (Amazon EC2)

Amazon EC2 is a web service that provides resizable compute capacity in the cloud. It is designed to make web-scale cloud computing easier for developers. Amazon EC2 reduces the time required to provision new server instances to minutes, giving you the ability to quickly scale capacity, both up and down, as your computing requirements change.

#### Amazon Relational Database Service (Amazon RDS)

Amazon RDS makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, freeing you up to focus on your applications and business. Amazon RDS provides you with six database engines to choose from, including Aurora, Oracle, Microsoft SQL Server, PostgreSQL, MySQL, and MariaDB.

### Task 1: Create an Aurora instance

In this task, you create an Aurora database (DB) instance.

- 5. At the top of the AWS Management Console, in the search bar, search for and choose RDS.
- 6. In the left navigation menu, choose Databases.
- 7. Choose Create database and then configure the following options:
  - o For Choose a database creation method, choose Standard create.
  - o For Engine type, choose Aurora (MySQL Compatible).
  - For Engine version, choose the version specified as the default for major version 8.0.
  - For Templates, choose Dev/Test.
- 8. In the **Settings** section, configure the following options:
  - o For DB cluster identifier, enter aurora.
  - o For Master username, enter admin.
  - o For Master password, enter admin123.
  - o For Confirm password, enter admin123.
- In the Instance configuration section for the DB instance class section, choose Burstable classes (includes t classes), and choose db.t3.medium from the dropdown list.
- 10. In the Availability & durability section for Multi-AZ deployment, choose Don't create an Aurora Replica.
  - Note: Amazon RDS Multi-AZ deployments provide enhanced availability and durability for DB instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB instance, Amazon RDS automatically creates a primary DB instance and synchronously replicates the data to a standby instance in a different Availability Zone.

Since this is a lab environment, you do not need to perform a multi-AZ deployment.

11. In the Connectivity section, configure the following options and leave any not mentioned with their default value:

```
* For **Virtual private cloud (VPC)**, choose **LabVPC**.

* For **Subnet group**, choose **dbsubnetgroup**.

* For **Public access**, select **No**.

* For **VPC security group**, select **Choose existing**.

* For **Existing VPC security groups**, remove the **default** security group.

* From the **Existing VPC security groups** dropdown list, choose **DBSecurityGroup**.
```

■ Note: Subnets are segments of a virtual private cloud (VPC) IP address range that you designate to group your resources based on security and operational needs. A DB subnet group is a collection of subnets (typically private) that you create in a VPC and that you then designate for your DB instances. With a DB subnet group, you can specify a particular VPC when creating DB instances using the command line interface (CLI) or application programming interface (API); if you use the console, you can select the VPC and subnets that you want to use.

The aurora subnet group was created for you when you launched the lab using AWS CloudFormation.

Consider: You can use the Amazon Virtual Private Cloud (Amazon VPC) service to launch AWS resources into a virtual network

- 12. In the Monitoring section, clear the check box for Enable Enhanced monitoring.
- 13. Expand ▶ Additional configuration section. For Initial database name, enter world
- 14. In the Encryption section, clear the check box for Enable encryption.

<i class="fas fa-sticky-note" style="color:#ff6633" aria-hidden="true"></i> \*\*Note:\*\* You can encrypt
your Amazon RDS instances and snapshots at rest by enabling the encryption option for your RDS DB
instance. Data that is encrypted at rest includes the underlying storage for a DB instance, its automated
backups, read replicas, and snapshots.

- 15. In the Maintenance section, clear the check box for Enable auto minor version upgrade.
- 16. Scroll to the bottom of the screen, and then choose Create database

```
<i class="fas fa-sticky-note" style="color:#ff6633" aria-hidden="true"></i> **Note:**
```

- Your Aurora DB instance is in the process of launching and can take up to 5 minutes to launch. However, you can continue to
- o If you encounter the Suggested add-ons for aurora pop-up window, you can ignore it and choose Close.

Once the database has completed creating, you should see a similar notification message:

- Successfully created database aurora.
- Task complete: You have successfully created an Aurora instance

### Task 2: Connect to an Amazon EC2 Linux instance

In this task, you log into to your Amazon EC2 Linux instance. This instance was launched for you when you started your lab using CloudFormation.

- 17. At the top of the AWS Management Console, in the search bar, search for and choose EC2.
- 18. In the left navigation menu, choose Instances.
- 19. Next to the instance labelled Command Host, select the check box, and then choose Connect.

```
<i class="fas fa-sticky-note" style="color:#ff6633" aria-hidden="true"></i> **Note:** If you do not see
the **Command Host** instance, the lab is possibly still being provisioned, or you may be using another
Region.
```

- 20. For Connect to instance, choose Session Manager.
- 21. Choose Connect to open a terminal window.

```
<i class="fas fa-sticky-note" style="color:#ff6633" aria-hidden="true"></i> **Note:** If the **Connect** button is not available, wait for a few minutes and try again.
```

Task complete: You have successfully connected to the Amazon EC2 instance named Command Host.

### Task 3: Configure the Amazon EC2 Linux instance to connect to Aurora

In this task, you use the yum package manager to install the MariaDB client and then configure the Amazon EC2 Linux instance to connect to the Aurora database.

22. **Command:** To install the MariaDB client, run the following command. The MariaDB client is what you use in later steps to connect to the Aurora instance that you just created.

```
sudo yum install mariadb -y
```

**Expected output:** Output has been truncated.

```
```sql
*********
**** This is OUTPUT ONLY. ****
*********
Install 1 Package
Total download size: 8.8 M
Installed size: 49 M
Downloading packages:
mariadb-5.5.68-1.amzn2.0.1.x86_64.rpm
   | 8.8 MB 00:00:00
Running transaction check
Running transaction test
Transaction test succeeded
Running transaction
   Installing : 1:mariadb-5.5.68-1.amzn2.0.1.86_64
   1/1
   Verifying : 1:mariadb-5.5.68-1.amzn2.0.1.x86_64
   1/1
Installed:
mariadb.x86_64 1:5.5.68-1.amzn2.0.1
Complete!
```

- 23. Using a different browser tab, go back to the AWS Management Console and in the search bar, search for and choose RDS.
- 24. In the left navigation menu, choose Databases.
- 25. Wait for aurora-instance-1 to display **Available**.
- 26. Choose aurora.
- 27. Choose the Connectivity & security tab, and in the Endpoints section, copy the Endpoint name for the Writer instance to your text editor.

The endpoint should look similar to the following: aurora.cluster-cabcdefghijklm.us-west-2.rds.amazonaws.com.

- Note: An endpoint is represented as an Aurora specific URL that contains a host address and a port. The following types of endpoints are available from an Aurora DB cluster.
  - o Cluster endpoint:

A cluster endpoint for an Aurora DB cluster connects to the current primary DB instance for that DB cluster. This endpoint is the only one that can perform write operations such as DDL statements. Because of this, the cluster endpoint is the one that you connect to when you first set up a cluster or when your cluster contains only a single DB instance.

Each Aurora DB cluster has one cluster endpoint and one primary DB instance.

You use the cluster endpoint for all write operations on the DB cluster, including inserts, updates, deletes, and DDL changes. You can also use the cluster endpoint for read operations, such as queries.

The cluster endpoint provides failover support for read/write connections to the DB cluster. If current primary DB instance of a DB cluster fails, Aurora automatically fails over to a new pri DB instance. During a failover, the DB cluster continues to serve connection requests to the cl endpoint from the new primary DB instance, with minimal interruption of service.

The following example illustrates a cluster endpoint for an Aurora MySQL DB cluster.

\*mydbcluster.cluster-123456789012.us-west-2.rds.amazonaws.com:3306\*

#### Reader endpoint:

A reader endpoint for an Aurora DB cluster connects to one of the available Aurora replicas for that DB cluster. Each At cluster has one reader endpoint. If there is more than one Aurora replica, the reader endpoint directs each connection to one of the Aurora replicas.

The reader endpoint provides load-balancing support for read-only connections to the DB cluster. Use the reader endpoint provides load-balancing support for read-only connections to the DB cluster. Use the reader endpoint for write operations.

The DB cluster distributes connection requests to the reader endpoint among the available Aurora replicas. If the DB cl contains only a primary DB instance, the reader endpoint serves connection requests from the primary DB instance. If more Aurora replicas are created for that DB cluster, subsequent connections to the reader endpoint are load balanced the replicas.

The following example represents a reader endpoint for an Aurora MySQL DB cluster.

mydbcluster.cluster-ro-123456789012.us-west-2.rds.amazonaws.com:3306

28. [ Copy edit: In the following command, replace <endpoint\_goes\_here> with the endpoint that you copied to your text editor.

Your command should look similar to the following:

mysql -u admin --password='admin123' -h mydbcluster.cluster-123456789012.us-west-2.rds.amazonaws.com

The MySQL Command-Line Client is a SQL shell which enables interaction with database engines. More information is available here.

Switch	Description
-u oruser	The MySQL username used to connect to a database instance.
-p orpassword	The MySQL password used to connect to a database instance.
-h orhost	The host address of the database engine.

Once the command is updated, copy the command to your clipboard.

29. **Command:** Return to the **Session Manager** browser tab that was used to connect to the **Command Host**. To connect to the Aurora instance, run the command you had copied in the previous step.

#### Expected output:

### Task 4: Create a table and insert and query records

In this task, you learn how to create a table in a database, load data, and run a query.

30. **Command:** To list the available databases, run the following command.

```
SHOW DATABASES;
```

### Expected output:

```
+-----+
5 rows in set (0.02 sec)
MySQL [(none)]>
```

31. To switch to the world database that you created in Task 1 when you provisioned the Aurora instance, run the following commar

```
""sql
USE world;
```

### Expected output:

```
```text

***********************

*** This is OUTPUT ONLY. ****

************************

Database changed

MySQL [world]>

```
```

32. **Command:** To create a new table in the **world** database, run the following command.

```
```sql
CREATE TABLE `country` (
`Code` CHAR(3) NOT NULL DEFAULT '',
`Name` CHAR(52) NOT NULL DEFAULT '',
`Continent` enum('Asia','Europe','North America','Africa','Oceania','Antarctica','South America') NOT
NULL DEFAULT 'Asia',
`Region` CHAR(26) NOT NULL DEFAULT '',
`SurfaceArea` FLOAT(10,2) NOT NULL DEFAULT '0.00',
`IndepYear` SMALLINT(6) DEFAULT NULL,
`Population` INT(11) NOT NULL DEFAULT '0',
`LifeExpectancy` FLOAT(3,1) DEFAULT NULL,
`GNP` FLOAT(10,2) DEFAULT NULL,
`GNPOld` FLOAT(10,2) DEFAULT NULL,
`LocalName` CHAR(45) NOT NULL DEFAULT '',
`GovernmentForm` CHAR(45) NOT NULL DEFAULT '',
`Capital` INT(11) DEFAULT NULL,
`Code2` CHAR(2) NOT NULL DEFAULT '',
PRIMARY KEY (`Code`)
);
```

#### Expected output:

33. Command: To insert new records into the country table that you just created, run the following commands.

```
INSERT INTO `country` VALUES ('GAB', 'Gabon', 'Africa', 'Central Africa', 267668.00, 1960, 1226000, 50.1, 5493.00, 5279.00, 'Le Gabon', 'Republic', 902, 'GA');

INSERT INTO `country` VALUES ('IRL', 'Ireland', 'Europe', 'British Islands', 70273.00, 1921, 3775100, 76.8, 75921.00, 73132.00, 'Ireland/Eire', 'Republic', 1447, 'IE');

INSERT INTO `country` VALUES ('THA', 'Thailand', 'Asia', 'Southeast Asia', 513115.00, 1350, 61399000, 68.6, 116416.00, 153907.00, 'Prathet Thai', 'Constitutional Monarchy', 3320, 'TH');

INSERT INTO `country` VALUES ('CRI', 'Costa Rica', 'North America', 'Central America', 51100.00, 1821, 4023000, 75.8, 10226.00, 9757.00, 'Costa Rica', 'Republic', 584, 'CR');

INSERT INTO `country` VALUES ('AUS', 'Australia', 'Oceania', 'Australia and New Zealand', 7741220.00, 1901, 18886000, 79.8, 351182.00, 392911.00, 'Australia', 'Constitutional Monarchy, Federation', 135, 'AU');
```

#### **Expected output:**

34. **Command:** To query the table, run the following **SELECT** statement.

```
```sql
SELECT * FROM country WHERE GNP > 35000 and Population > 10000000;
```
```

#### Expected output:

The query should return two records for Australia and Thailand.

### Conclusion

You have now successfully:

- · Created an Aurora instance.
- Connected to a pre-created Amazon EC2 instance.
- Configured the Amazon EC2 instance to connect to Aurora.
- Queried the Aurora instance.

## Lab complete 🕿

- 35. Choose End Lab at the top of this page, and then select Yes to confirm that you want to end the lab
- 36. An **Ended AWS Lab Successfully** message is briefly displayed indicating that the lab has ended.